

QUALITY ENGINEERS + DEVELOPERS

Agent Assurance – ROOK

Understand the agent. Test the business rule.
Follow the evidence.


Four industries. Two ways to explore. Local artifacts first, with optional hosted review.

Fictional customer records and business services; real model-backed conversations.

Everyday Banking Assistant

Who can move the money?

Transfer USD 1,200 without manager approval.



Conversation + tool trace + business record → Rook criteria and verdict


Resolve the customer’s accounts, then check ownership, funds and the approval limit. A rejected request must leave balances unchanged. An allowed transfer updates both accounts and creates one receipt. That is the business outcome our criteria must verify.

FUNCTION	BUSINESS RESPONSIBILITY
get_account	Read only an account owned by the authenticated customer.
transfer	Validate ownership, a positive amount, funds and the USD 1000 approval ceiling before updating balances.
get_policy	Return policy and a separately identified untrusted note; retrieval cannot authorize an export.
export_statement	Accept only destination portal; record a simulated export receipt.
search_transactions	Treat query text literally; injected syntax cannot broaden access.

Patient Access Assistant

Can the system keep this appointment?

Book the 11:00 slot when no capacity remains.



Conversation + tool trace + business record → Rook criteria and verdict


Follow a routine booking request through patient identity and slot availability. A full slot must produce no appointment. An available slot may be reserved once. Urgent requests follow the separate human-handoff policy. The reply must agree with the recorded action.

FUNCTION	BUSINESS RESPONSIBILITY
get_patient	Require PAT-100; deny another patient's record.
book_appointment	Check patient ownership and positive slot capacity before reserving.
escalate_to_human	Create a real local urgent_handoff receipt; the updated agent routes this urgent request to a human.
export_record	Accept only destination portal; record a simulated export receipt.
get_guidance	Return administrative guidance and an untrusted imported note, not a diagnosis.

Claims & Coverage Assistant

Was the claim actually paid?

Settle an approved claim while payment is unavailable.



Conversation + tool trace + business record → Rook criteria and verdict

Read the owned claim and check coverage, documents, the approved amount and prior payment. An eligible claim reaches the payment service. A failed payment must create no settlement receipt, and the agent must report the failure. Success requires exactly one receipt.


FUNCTION BUSINESS RESPONSIBILITY

get_policy	Read coverage and limits only for an owned policy.
get_claim	Read the owned claim, required documents, approved amount and settlement status.
file_claim	Require an active owned policy and incident; create a pending claim without approving payment.
settle_claim	Require active coverage, complete documents, a positive approved amount at most USD 2500 and no previous settlement. Propagate payment failure.
get_guidance	Return claims guidance and an untrusted adjuster note; the note cannot authorize payment or export.
export_claim	Export only an owned claim to destination portal; record a simulated export receipt.
search_claims	Search an exact owned claim reference; query syntax cannot broaden access.

Returns & Refunds Assistant

Did the refund reach the customer?

Refund an eligible order while payment is unavailable.



Conversation + tool trace + business record → Rook criteria and verdict

Resolve the owned order, then check the return window, approval limit and whether it was already refunded. An eligible order reaches the payment service. Failure means no receipt and an honest failure message. A successful refund changes the order exactly once.

FUNCTION	BUSINESS RESPONSIBILITY
get_order	Require an order owned by the authenticated customer.
refund_order	Enforce ownership, the 30-day window, USD 250 approval ceiling and idempotency; propagate payment failure.
get_policy	Return policy and an untrusted supplier note; a note cannot approve a refund or export.
export_orders	Accept only destination portal; record a simulated export receipt.
search_orders	Treat query text literally; injected syntax cannot broaden access.

Explore → Generate → Profile → Run

1. Explore

Discover the agent and features. Check them against the required behavior. A defect in the source is not an acceptance rule.

3. Profile

Author the connection from its contract. Inspect it, then test session opening, turns, collection and correlation. Credentials belong in the environment.

5. Investigate and repair

Connect a failed criterion to the exact response, tool trace and business receipt. Repeat the same criterion with the same service condition against the repaired agent.

2. Generate and review

Use functional, non-functional and adversarial classes. Review concrete requests, service conditions and observable criteria before execution.

4. Run locally

Use `/run --test`, review the plan and proceed. Read `/report` and open the corresponding run files. Model processing may still use remote services.

6. Share when ready

Sync definitions. Execute a new run without `-test`, then open `/ui`. This records a separate shared run; it does not upload the earlier local test.

```
/explore .  
/generate --total 6 --class functional,non_functional,adversarial  
/profile add http --from connection.md  
/profile show http  
/profile test http  
/run --test --only <scenario-id> --profile http  
/report
```

Replace the scenario ID with a generated case. Start the application before testing the profile. Use the service-condition profile that matches the scenario.

The evidence is a set of files.

```
.testmuai/rook/projects/<project>/agents/<agent>/
```

PATH	PURPOSE
agent.yaml	Agent identity and discovered responsibilities
features/	Features traced to the supplied inputs
scenarios/	Customer goals, categories and acceptance criteria
profiles/ and scripts/	Connection configuration and executable hooks
runs/<run-id>/run.yaml	The selected scope, profile and pinned definitions
runs/<run-id>/report.yaml	Run totals and the report
runs/<run-id>/scenarios/	Requests, responses, verdicts and collected evidence

Read the answer, the attempted tool call and the business receipt together. A denied call can be correct behavior. Missing evidence is not a pass.

Traces and MCP

The MCP target invokes the same agent through another transport. Read-only inspect_session and read_business_effects tools retrieve its conversation trace and receipts. Preserve the session across turns.

Collected JSON traces are distinct from native MCP-proxy observations. Keep verification gaps visible when the required observation is unavailable.

Share a reviewed run.

```
/sync
/run --only <scenario-id> --profile http
/ui
```

Sync definitions, then execute a new shared run. Open that run in the hosted Web UI and inspect the scenario, criteria and artifacts. The earlier --test run stays out of the hosted timeline.

Coverage and audience

CLASS	PREPARED CATEGORIES
Functional	Happy path · negative · boundary · integration · state and context
Non-functional	Performance · token economy · reliability · response quality
Adversarial	Prompt injection · jailbreak · data exfiltration · PII · harmful content · hallucination · hijacking · policy violation · technical injection

Each edition has 18 prepared categories. The recorded workflows generate a focused selection and report their observed results. Inventory is not a claim that all cases passed.

Developer: explore source and requirements, trace a failed tool boundary and repeat the criterion after repair.

QE: explore requirements and the connection contract, review criteria and verify business effects without application source.