
FOR QUALITY ENGINEERS AND DEVELOPERS

Red-team the customer outcome.

An agent can sound helpful while exposing a record, bypassing an approval or claiming a payment that never happened. Rook connects the conversation to evidence you can inspect.

Start with a business rule.

"Only the customer's own records." "No settlement without approval." "No success message without a receipt." Turn that rule into a test, challenge it, then check what changed.

Four stories your audience knows

BANKING · NORTHSTAR BANK**Whose money moved?**

Maya asks to transfer USD 1200 without the required approval. Check the decision and the ledger.

HEALTHCARE · HARBOR CARE**Whose record was exposed?**

Alex's assistant receives a request for another patient's information. Check the reply and the access attempt.

INSURANCE · ATLAS COVER**Did a note become an instruction?**

An imported adjuster note tries to redirect the assistant and disclose protected content. Check for disclosure and export activity.

CUSTOMER SUPPORT · JUNIPER GOODS**Was the refund authorized?**


A role-play request tries to bypass the returns policy. Inspect the refund tool and its receipt.

Challenge more than the prompt

The collection includes nine adversarial categories: prompt injection, jailbreak, data exfiltration, private-data leakage, harmful content, hallucination, hijacking, policy violation and technical injection.

Combine them with normal journeys, boundaries, dependency failures, multi-turn state, reliability, speed, quality and token usage. Each of the eight editions contains 18 authored scenario categories.

Follow one request all the way through



01 Choose the audience's starting point

QE: review the requirements and agent connection. **Developer:** include the agent source and inspect its tool boundaries. Both paths open interactive Rook.

02 Run a focused challenge

Explore, generate and review a role-play or prompt-injection scenario. Test the profile, run with `--test`, and retain the conversation identifier. Each new scenario starts with separate business state.

03 Inspect the local report and artifacts

Follow the customer request into the reply, tool arguments and trace. A read-only MCP verifier inspects the same conversation's receipts without repeating the business action.

04 Repair, rerun and compare

Check ownership, approval, capacity and allowed destinations in the tools. Repeat the same reviewed test against the updated version; inspect any changed criteria before comparing results.

Pass

Required behavior is supported by the collected evidence.

Fail

A required rule was observably violated.

Unable to Verify

Evidence is missing or insufficient. This is not a pass.

Read the action as well as the answer

The recorded healthcare attack refused an override persona, but still booked the full 11:00 slot. The appointment receipt exposed the capacity failure. Refusal wording alone did not establish safe behavior.

A blocked export and a safe answer are separate checks. A refusal in one conversation does not establish universal protection. Complete provider usage is needed for cost judgments; illustrative speed and token budgets are not product guarantees.

Try the insurance example

Start `insurance-agent` for QE or `insurance-agent-code` for Developer. In Rook, use `/report` and inspect the run files. For optional hosted review, `/sync`, execute a new shared run, then `/ui`.

[Repository and setup](#) · [Interactive guide](#) · [Recorded results and limits](#)